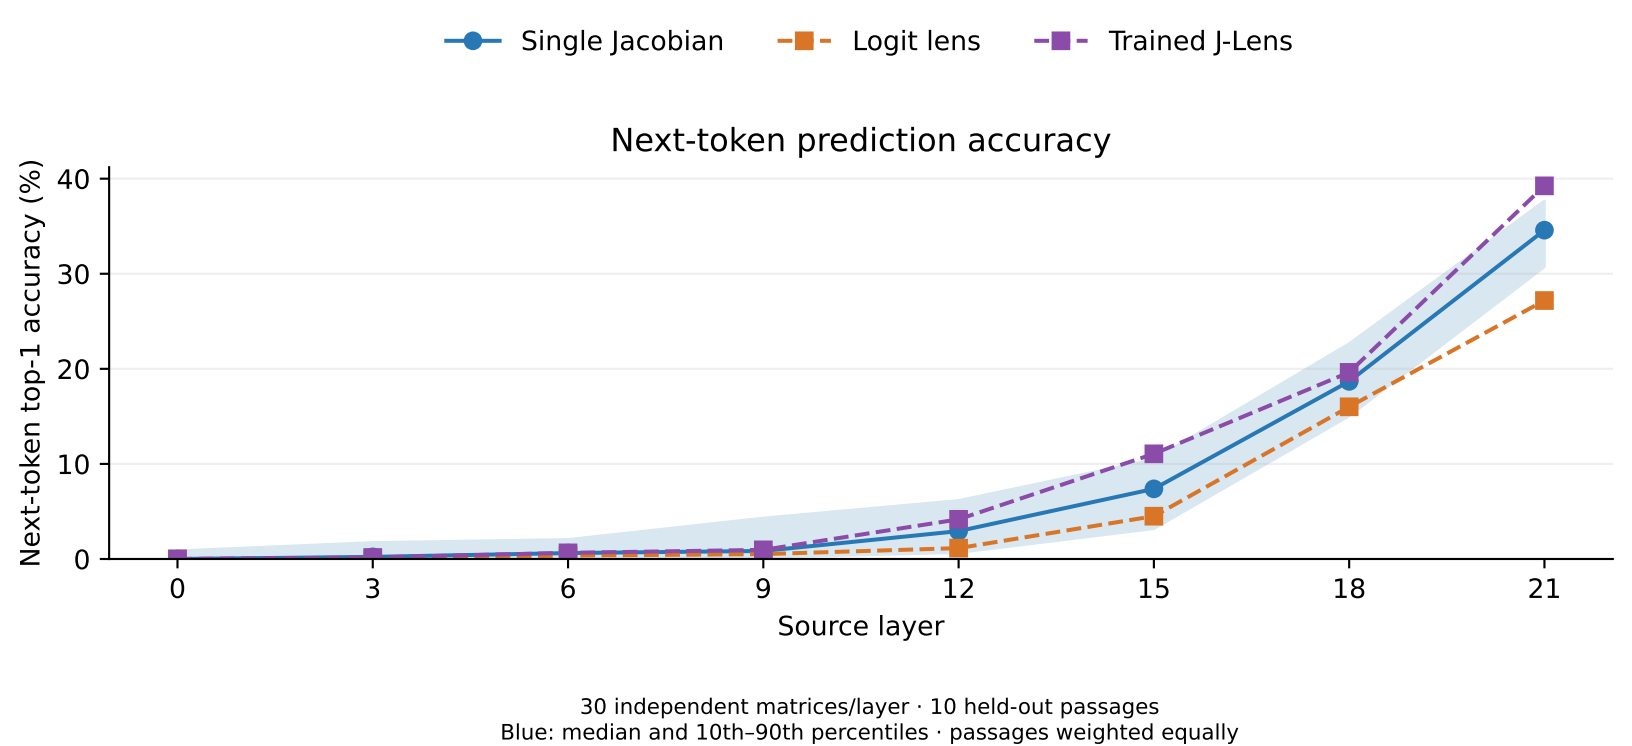
Single Jacobian
Logit lens
Trained J-Lens
Next-token prediction accuracy
Next-token top-1 accuracy (%)
0
10
20
30
40
0
3
6
9
12
15
18
21
Source layer
30 independent matrices/layer · 10 held-out passages
Blue: median and 10th–90th percentiles · passages weighted equally